

✓ Model1_TL - Transfer learning using InceptionNET

by Kaushik Srivatsan - CDS - kaushik.s-25@scds.saiuniversity.edu.in

- Trained on Hutton Rock Dataset
- Model contains BatchNormalization and Dropout of 0.25% Drop Rate before output layer

```
#importing the file
import zipfile
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
# zip_ref = zipfile.ZipFile("/content/drive/MyDrive/DL_Project/Hutton_Rock.zip", 'r')
# zip_ref.extractall("/content/drive/MyDrive/DL_Project/")
# zip_ref.close()
```

```
#Importing the libraries
import tensorflow as tf
print(tf.__version__)

from tensorflow import keras
tf.random.set_seed(42)

import numpy as np
np.random.seed(42)

import matplotlib.pyplot as plt
%matplotlib inline

import glob
import PIL
from PIL import Image
```

2.15.0

```
#Importing the images using glob as ImgFiles
imgFiles = glob.glob("/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/*.jpg")
for items in imgFiles[:8]:
    print(items)
```

```
/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/granite/1.jpg
/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/granite/125.jpg
/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/granite/106.jpg
/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/granite/117.jpg
/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/granite/127.jpg
/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/granite/116.jpg
/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/granite/10.jpg
/content/drive/MyDrive/DL_Main/DL_Project/Hutton_Rock/granite/12.jpg
```

✓ Data Preprocessing and Labelling

```
X = []
y = []

for fName in imgFiles:

    # Prepare the dataset and populate X and y
    X_i = Image.open(fName)
    X_i = X_i.resize((299,299))
    X_i = (np.array(X_i).astype(np.float32) /127.5) - 1

    if X_i.shape == (299, 299, 3):
        X.append(X_i)

        label = fName.split("/")
        y_i = label[-2]

        y.append(y_i)

print(y)
```

```
[[[ 0.27843142  0.2941177  0.254902 ]
 [ 0.27843142  0.2941177  0.254902 ]
 [ 0.27058828  0.28627455  0.24705887]
 ...
 [ 0.12941182  0.12941182  0.12941182]
 [ 0.04313731  0.04313731  0.04313731]
 [ 0.01176476  0.01176476  0.01176476]]]
```

```

[[ 0.27843142 0.2941177 0.254902 ]
 [ 0.27843142 0.2941177 0.254902 ]
 [ 0.27843142 0.2941177 0.254902 ]
 ...
 [ 0.16078436 0.16078436 0.16078436]
 [ 0.082353 0.082353 0.082353 ]
 [ 0.05098045 0.05098045 0.05098045]]

[[ 0.27843142 0.2941177 0.254902 ]
 [ 0.28627455 0.30196083 0.26274514]
 [ 0.2941177 0.30980396 0.27058828]
 ...
 [ 0.23921573 0.2313726 0.23921573]
 [ 0.17647064 0.17647064 0.17647064]
 [ 0.15294123 0.15294123 0.15294123]]

...

[[ 0.07450986 0.02745104 0.02745104]
 [ 0.082353 0.03529418 0.03529418]
 [ 0.09019613 0.04313731 0.04313731]
 ...
 [-0.04313725 -0.08235294 -0.05098039]
 [-0.09019607 -0.12941176 -0.09803921]
 [-0.10588235 -0.14509803 -0.11372548]]

[[ 0.082353 0.03529418 0.03529418]
 [ 0.09019613 0.04313731 0.04313731]
 [ 0.12156868 0.07450986 0.07450986]
 ...
 [-0.05882353 -0.09803921 -0.06666666]
 [-0.10588235 -0.14509803 -0.11372548]
 [-0.12156862 -0.1607843 -0.12941176]]

[[ 0.082353 0.03529418 0.03529418]
 [ 0.09803927 0.05098045 0.05098045]
 [ 0.13725495 0.09019613 0.09019613]
 ...
 [-0.06666666 -0.10588235 -0.0745098 ]
 [-0.11372548 -0.15294117 -0.12156862]
 [-0.12941176 -0.16862744 -0.1372549 ]]]

[[[-0.09019607 -0.02745098 -0.02745098]
 [ 0.11372554 0.18431377 0.21568632]
 [ 0.06666672 0.12156868 0.18431377]
 ...
 [-0.1372549 -0.1372549 -0.0745098 ]
 [-0.05882353 -0.0745098 -0.0745098 ]
 [-0.11372548 -0.1607843 -0.17647058]]

```

✓ Transfer Learning - Using InceptionNet - Importing InceptionResNetV2

```
base_inceptionnet = keras.applications.InceptionResNetV2(weights='imagenet', input_shape = (299,299,3), include_top=False)
```

```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/inception\_resnet\_v2/inception\_resnet\_v2\_weights\_1219055592/219055592 [=====] - 1s 0us/step

```

```
base_inceptionnet.trainable = False
```

```

for layer in base_inceptionnet.layers:
    layer.trainable = False

```

✓ Building the Classifier

```

global_pool = keras.layers.GlobalAveragePooling2D()(base_inceptionnet.output)
x = keras.layers.BatchNormalization()(global_pool)
x = keras.layers.Dropout(0.25)(x)

output_ = keras.layers.Dense(units=4, activation='softmax')(x)
model_TL1 = keras.models.Model(inputs=[base_inceptionnet.input], outputs=[output_])

```

✓ Training the Model

```

model_TL1.compile(loss='sparse_categorical_crossentropy', optimizer='adam',metrics=['accuracy'])

callbacks_TL = [keras.callbacks.ModelCheckpoint("bestM1TL.h5",monitor='val_accuracy',save_weights_only=True,save_best_only=True)]

history_TL = model_TL1.fit(x = X_train, y = y_train, epochs=10, batch_size = 16, validation_split=0.1, callbacks = callbacks_TL)

Epoch 1/10
21/21 [=====] - 38s 838ms/step - loss: 1.4030 - accuracy: 0.4806 - val_loss: 1.0219 - val_accuracy: 0.4737
Epoch 2/10
21/21 [=====] - 5s 238ms/step - loss: 0.6894 - accuracy: 0.7522 - val_loss: 0.8851 - val_accuracy: 0.6316
Epoch 3/10
21/21 [=====] - 4s 186ms/step - loss: 0.6043 - accuracy: 0.7672 - val_loss: 0.8194 - val_accuracy: 0.6053
Epoch 4/10
21/21 [=====] - 5s 247ms/step - loss: 0.4828 - accuracy: 0.8090 - val_loss: 0.7703 - val_accuracy: 0.6579
Epoch 5/10
21/21 [=====] - 4s 181ms/step - loss: 0.4157 - accuracy: 0.8537 - val_loss: 0.6672 - val_accuracy: 0.6579
Epoch 6/10
21/21 [=====] - 4s 182ms/step - loss: 0.4164 - accuracy: 0.8448 - val_loss: 0.6412 - val_accuracy: 0.6316
Epoch 7/10
21/21 [=====] - 5s 263ms/step - loss: 0.3436 - accuracy: 0.8567 - val_loss: 0.5925 - val_accuracy: 0.6842
Epoch 8/10
21/21 [=====] - 5s 254ms/step - loss: 0.3283 - accuracy: 0.8806 - val_loss: 0.5628 - val_accuracy: 0.7368
Epoch 9/10
21/21 [=====] - 4s 184ms/step - loss: 0.3161 - accuracy: 0.8776 - val_loss: 0.5490 - val_accuracy: 0.6842
Epoch 10/10
21/21 [=====] - 4s 207ms/step - loss: 0.2881 - accuracy: 0.9015 - val_loss: 0.5355 - val_accuracy: 0.7368

```

✓ Testing the Model using Best Weights

```

model_TL1.load_weights("bestM1TL.h5")

testLoss_TL, testAccuracy_TL = model_TL1.evaluate(x = X_test, y = y_test)

print("Test-loss: %f, Test-accuracy: %f" % (testLoss_TL, testAccuracy_TL))

4/4 [=====] - 2s 318ms/step - loss: 0.7360 - accuracy: 0.7200
Test-loss: 0.736017, Test-accuracy: 0.720000

```

✓ Checking the performance

```

y_proba = model_TL.predict(X_test_std)
y_predict = np.argmax(y_proba, axis=-1)

4/4 [=====] - 12s 2s/step

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_true = y_test, y_pred = y_predict)

fig, ax = plt.subplots(figsize=(6, 6))
ax.matshow(cm, cmap=plt.cm.Blues, alpha=0.3)


for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(x=j, y=i, s=cm[i, j], va='center', ha='center')

ax.title.set_text('CNN\n')
plt.xlabel('Predicted label')
plt.ylabel('True label')

plt.tight_layout()
plt.savefig("ConfusionMatrix.png", dpi=300, format='png', pad_inches=0.3)
plt.show()

print(set(y))
print(lEncoder.classes_)

```



▼ Saving the Model

```
model_TL1.save('/content/drive/MyDrive/Models/Model1_TL.h5')
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103: UserWarning: You are saving your model as an HDF5 file via the save method. This file format is deprecated, and we recommend you instead use the save_model method, which saves files in the modern Keras format.
  saving_api.save_model(
```

```
from numpy import save
```

```
save('/content/drive/MyDrive/Models/X_train_std_model1.npy', X_train_std)
save('/content/drive/MyDrive/Models/X_test_std_model1.npy', X_test_std)
```

```
save('/content/drive/MyDrive/Models/y_train_model1.npy', y_train)
save('/content/drive/MyDrive/Models/y_test_model1.npy', y_test)
```